

TEJ GANGUPANTULA

+1 (209)-600-5060 | tejgangupantula@gmail.com | [Portfolio](#) | [LinkedIn](#) | [GitHub](#)

EDUCATION

University of California, Santa Barbara

September 2023 – December 2026

B.S. Statistics and Data Science

Santa Barbara, CA

- **GPA:** 3.6
- **Relevant Courses:** Data Science with R/SQL, Probability and Statistics, Data Science Concepts and Analysis, Regression Analysis, Statistical Machine Learning, Stochastic Models & Applied Probability, Calculus I-III, Linear Algebra, Differential Equations

SKILLS

Languages: Python, R, SQL, Javascript, Typescript, Java

Libraries: Scikit-learn, LangGraph, LangChain, pandas, NumPy, Matplotlib, Seaborn, TensorFlow, spaCy, Streamlit

Tools/Frameworks: Git, Jupyter, Google Cloud Platform, FastAPI, PostgreSQL, FastMCP, Docker

PROFESSIONAL EXPERIENCE

AI Engineer Intern

Summer 2025, Spring & Summer 2026

Agentman

Berkeley, CA

- Developed production-ready MCP (Model Context Protocol) servers in TypeScript/Javascript, containerized with Docker and deployed to GCP Cloud Run, enabling agents to interact with Gmail (read, compose, edit) and an image generation service (Nano Banana) as callable tools.
- Built and integrated agentic AI workflows for healthcare operations using HIPAA-compliant APIs and GCP, connecting EHR data sources and automating clinical administrative tasks, reducing manual processing by ~70% and saving practices \$20K+ per user annually.
- Designed cloud pipelines connecting third-party APIs to agent tool layers, streamlining operational workflows and saving 3+ staff-hours daily with improved reliability through schema validation and QA.

Machine Learning Intern

July 2024 – September 2024

FusionCare

Davis, CA

- Built a Python-based clinical text processing pipeline to ingest, clean, and structure de-identified patient notes, applying NLP-based PII masking (spaCy/medspaCy) to safely prepare data for downstream analysis.
- Integrated OpenAI GPT models via LangChain with structured prompting and output validation to convert unstructured clinical notes into JSON-formatted visit summaries, improving care plan evaluations and saving clinicians 4–5 hours weekly.
- Prototyped and iterated on outputs using Streamlit, enabling rapid feedback loops with clinical stakeholders during development.

PROJECTS

Prior Authorization Voice Agent | *Python, LangGraph, LangChain, FastAPI, PostgreSQL*

April 2026 - June 2026

- Built a real-time voice agent (STT → LangGraph state machine → TTS) that answers prior-authorization status questions over a browser interface, routing callers across four normalized outcomes via LangGraph's conditional edges rather than an open-ended ReAct loop.
- Integrated the ElevenLabs API for real-time speech-to-text and text-to-speech, and designed three consent-gated tools (PA lookup, appeal ticket, escalation ticket) backed by a PostgresSaver checkpoint for durable per-call state that survives process restarts.
- Authored an architecture-first design that pairs every demo shortcut with its production successor, and implemented structured error handling with per-tool routing and a dedicated escalation path for total failures.

NBA Finals Outcome Prediction Analysis | *Python, Scikit-learn, NumPy, pandas*

April 2025 – June 2025

- Developed a Random Forest classification model in Python to simulate and predict 2025 NBA Finals outcomes, achieving 62% test accuracy using team-level regular season and playoff metrics.
- Engineered a Monte Carlo simulation to estimate championship probabilities over a best-of-seven series, predicting a 75.5% chance of the Oklahoma City Thunder winning against the Indiana Pacers (24.5%).
- Cleaned, processed, and analyzed historical NBA data with pandas, numpy, and scikit-learn, enabling interpretable model outputs and supporting data-driven sports analytics.